
Table of Contents

rungreedybandit.m	1
Inputs:	1
Outputs:	2
Code:	2

rungreedybandit.m

```
% Runs Greedy Bandit algorithm and returns regret and fraction of
% pulls.
%
% This code implements the Greedy Bandit algorithm. For more
% information,
% see our paper
%
% - https://arxiv.org/abs/1704.09011.
%
% The algorithm proceeds as follows:
% First each arm is forced sampled by some number of time periods
% given by random_initialization). Then, at each time period the arm
% with
% the highest estimated mean is pulled. The algorithm uses Sherman-
% Morrison
% formula for fast updates of least squares estimations of arms. This
% fast update is applied once the covariance matrices become
% invertible.
% Before having invertible covariance matrices, the algorithm applies
% ridge
% regression for more accurate estimations.
%
```

Inputs:

k: Number of **arms**.
T: Time horizon.
d: Dimension of **covariates**.
b: A k*d matrix of **arm parameters**.
sigma_e: Standard deviation of **noise** (used only for **reward generation**).
 This **parameter is unused if noise and contexts are provided**.
sigma_x: Standard deviation of **covariates**
 (used only **for context generation for Gaussian contexts**).
 This **parameter is unused if noise and contexts are provided**.
xmax: Maximum of **l2-norm of covariates**
 (used only **for context generation**). This **parameter is unused if noise and contexts are provided**.
random_initialization: Number of **rounds of forced sampling per arm in the beginning**. If set to zero, the algorithm is purely greedy.
verbose: Whether to **print outputs or not**.

varargin: Additional arguments. In particular, if these are not provided the noise and contexts will be generated according to Gaussian and truncated Gaussian distributions. In case they are provided, there should exactly be THREE additional arguments. The first one is contexts. The second one is a binary input, called noise_input. If noise_input = 1, this means the last argument will be noise $e=(Y-X\beta)$. On the other hand, if noise_input = 0, then the last argument will be Y or rewards. Note that the noise should be $T \times 1$ while rewards should be $T \times k$.

Outputs:

regret: Cumulative regret as a running sum over regret terms.
fractions: Fractions of pulls of different arms.

Code:

```
function [regret, fractions] = rungreedybandit(k, T, d, b,
    sigma_e, ...
    sigma_x, xmax, random_initialization, verbose, varargin)

warning('off','all');

if nargin==9 % Context and noise are NOT provided, so generate those.
    % Noise is Gaussian with std sigma_e.
    e = randn(T,1)*sigma_e;
    noise_input = 1;

    % Contexts follow truncated gaussian distributions with l-infinity
    norm
    % at most xmax.
    X = max(-xmax, min(xmax, mvnrnd(zeros(d, 1), sigma_x, T)));
else
    X = varargin{1};
    noise_input = varargin{2};
    if(noise_input==1)
        e = varargin{3};
    else
        rewards = varargin{3};
    end
end

reward_vector = zeros(T, k); % Vector of all (potential) rewards.
pull_ind = zeros(T, k); % Binary indicator whether each is pulled.

regret = zeros(1, T);
betahat = b * 0; % Initialize all arm estimations with vector of
    zeros.

XtopX_inv = zeros(d, d, k);
```

```

% Whether direct LS calculation using normal equations is needed.
no_LS_calculations = zeros(k, 1);

for t=1:T
    x = X(t,:);

    %----- First, choose which arm to pull this round.

    if (t>random_initialization * k)
        z = betahat * x;
        opt_arms = find(z==max(z));
        % Break ties randomly.
        arm_pulled = opt_arms(randi(length(opt_arms)));
    else
        arm_pulled = mod(t-1, k) + 1;
    end
    pull_ind(t, arm_pulled) = 1;

    %----- Second, calculate the regret.

    if(noise_input==1)
        bx = b*x;
        ourreward = bx(arm_pulled);
        bestreward = max(bx);
    else
        ourreward = rewards(t, arm_pulled);
        bestreward = max(rewards(t,:));
    end

    if (t==1)
        regret(t) = bestreward - ourreward;
    else
        regret(t) = regret(t-1) + bestreward - ourreward;
    end

    %----- Third, update estimates.

    if(noise_input==1)
        reward_vector(t, arm_pulled) = ourreward + e(t);
    else
        reward_vector(t, arm_pulled) = rewards(t, arm_pulled);
    end

    if (no_LS_calculations(arm_pulled)==0)
        obs_filt = find(pull_ind(:, arm_pulled)==1); % Filter
        lsX = X(obs_filt, :); % Design matrix.
        lsY = reward_vector(obs_filt, arm_pulled); % Observations.
        if (rank(lsX)>=d)
            XtopX_inv(:, :, arm_pulled) = inv(lsX'*lsX);
            betahat(arm_pulled, :) = lsX\lsY; %
            no_LS_calculations(arm_pulled) = 1;
        else

```

```

        % Empirical estimator of  $\|x\|_2$ , used for ridge
regression.
        penalty = norm(lsX)/size(lsX,1);
        hat_beta = (lsX'*lsX + penalty^2 * eye(d)) \ (lsX' * lsY);
        betahat(arm_pulled,:) = hat_beta';
    end
    else
        [XtopX_inv(:, :, arm_pulled), betahat_vertical] =
rankoneupdate(...
        XtopX_inv(:, :, arm_pulled), betahat(arm_pulled, :)',
x, ...
        reward_vector(t, arm_pulled));
        betahat(arm_pulled,:) = betahat_vertical';
    end
    if (verbose==1)
        if (mod(t,500)==0)
            fprintf('GB: t=%d, parameter estimation error = %f. \n',
t, ...
                    norm(b - betahat, 'fro'))
        end
    end
end
fractions = mean(pull_ind); % Fraction of times each arm is pulled.
if(verbose==1)
    fprintf('GB: Total parameter estimation error = %f. \n', ...
            norm(b - betahat, 'fro'));
    fprintf('GB: Fraction of pulls = %f. \n', fractions);
    fprintf('GB: Total regret occurred = %f. \n', regret(end));
end
end

```

Published with MATLAB® R2015a