
Table of Contents

rungreedyfirst.m	1
Inputs:	1
Outputs:	2
Code:	2

rungreedyfirst.m

```
% Runs Greedy First algorithm and returns regret and fraction of
% pulls.
%
% This code implements the Greedy First algorithm. For more
% information,
% see our paper:
%
% - https://arxiv.org/abs/1704.09011.
%
% The algorithm proceeds as follows: first each arm is forced sampled
% by
% some number of time periods given by random_initialization). Then,
% at
% each time period the arm with the highest estimated mean is pulled.
% The algorithm keeps track of the %smallest value among minimum
% eigenvalues of the covariance matrices of arms and switches to OLS
% Bandit
% algorithm if this number drops below a linearly growing (in t)
% threshold.
% For more details on OLS bandit and discussion about its parameters
% see:
%
% - https://pubsonline.informs.org/doi/abs/10.1287/11-SSY032.
% - https://pubsonline.informs.org/doi/10.1287/opre.2019.1902.
%
% The algorithm uses Sherman-Morrison formula for fast updates of
% least
% squares estimations of arms. This fast update is applied once the
% covariance matrices become invertible. Before having invertible
% covariance matrices, the algorithm applies ridge regression for more
% accurate estimations.
%
```

Inputs:

```
k: Number of arms.
T: Time horizon.
d: Dimension of covariates.
b: A k*d matrix of arm parameters.
sigma_e: Standard deviation of noise (used only for reward
generation).
```

This parameter is unused if noise and contexts are provided.
sigma_x: Standard deviation of covariates
(used only for context generation for Gaussian contexts).
This parameter is unused if noise and contexts are provided.
xmax: Maximum of l2-norm of covariates
(used only for context generation). This parameter is unused if noise and contexts are provided.
h: h parameter used in OLS bandit algorithm.
q: q parameter used in OLS bandit algorithm.
t0: Time at which the algorithm starts checking minimum eigenvalues from then onwards.
min_eig_threshold: Bound on the growth of smallest minimum eigenvalue among covariance matrices.
random_initialization: Number of rounds of forced sampling per arm in the beginning. If set to zero, the algorithm is purely greedy.
verbose: Whether to print outputs or not.
varargin: Additional arguments. In particular, if these are not provided the noise and contexts will be generated according to Gaussian and truncated Gaussian distributions.
In case they are provided,
there should exactly be THREE additional arguments. The first one is contexts. The second one is a binary input, called noise_input. If noise_input = 1, this means the last argument will be noise $e=(Y-X\beta)$. On the other hand, if noise_input = 0, then the last argument will be Y or rewards. Note that the noise should be $T \times 1$ while rewards should be $T \times k$.

Outputs:

regret: Cumulative regret as a running sum over regret terms.
fractions: Fractions of pulls of different arms.
switched: The time at which Greedy-First switches to OLS bandit. If the algorithm does not switch, this is equal to -1.

Code:

```
function [regret, fractions, switched] = rungreedyfirst(k, T, d,
    b, ...
    sigma_e, sigma_x, xmax, h, ...
    q, t0, min_eig_threshold, random_initialization, ...
    verbose, varargin)

warning('off','all');

switched = -1; % This only will be set if algorithm switches from
greedy.

if nargin==13 % Context and noise are NOT provided, so generate those.
    % Noise is Gaussian with std sigma_e.
    e = randn(T, 1)*sigma_e;
    noise_input = 1;
```

```

    % Contexts follow truncated gaussian distributions with l-infinity
    norm
    % at most xmax.
    X = max(-xmax, min(xmax, mvnrnd(zeros(d, 1), sigma_x, T)));
else
    X = varargin{1};
    noise_input = varargin{2};
    if(noise_input==1)
        e = varargin{3};
    else
        rewards = varargin{3};
    end
end

reward_vector = zeros(T, k); % Vector of all (potential) rewards.
pull_ind = zeros(T, k); % Binary indicator whether each is pulled.

regret = zeros(1, T);
betahat = b * 0; % Initialize all arm estimations with vector of
    zeros.

XtopX_inv = zeros(d, d, k);

% Whether direct LS calculation using normal equations is needed.
no_LS_calculations = zeros(k, 1);

% Minimum eigenvalue of sample covariance matrix of different arms.
lam_hat = zeros(k, 1);
lam0 = inf;

switch_to_olsbandit = 0; % Binary, whether algorithm switches or not.

t = 1;

while (t<=T && switch_to_olsbandit==0)
    x = X(t,:)' ;
    %----- First, choose which arm to pull this round.
    if (t>random_initialization * k)
        z = betahat * x;
        opt_arms = find(z==max(z));
        % Break ties randomly.
        arm_pulled = opt_arms(randi(length(opt_arms)));
    else
        arm_pulled = mod(t-1, k) + 1;
    end
    pull_ind(t, arm_pulled) = 1;

    %----- Second, calculate the regret.

    if(noise_input==1)
        bx = b*x;
        ourreward = bx(arm_pulled);
        bestreward = max(bx);

```

```

else
    ourreward = rewards(t, arm_pulled);
    bestreward = max(rewards(t,:));
end

if (t==1)
    regret(t) = bestreward - ourreward;
else
    regret(t) = regret(t-1) + bestreward - ourreward;
end

%----- Third, update estimates.

if(noise_input==1)
    reward_vector(t, arm_pulled) = ourreward + e(t);
else
    reward_vector(t, arm_pulled) = rewards(t, arm_pulled);
end

if (no_LS_calculations(arm_pulled)==0)
    obs_filt = find(pull_ind(:, arm_pulled)==1); % Filter
observations.
    lsX = X(obs_filt, :); % Design matrix.
    lsY = reward_vector(obs_filt, arm_pulled); % Observations.
    if (rank(lsX)>=d)
        XtopX_inv(:, :, arm_pulled) = inv(lsX'*lsX);
        betahat(arm_pulled, :) = lsX\lsY; %
        no_LS_calculations(arm_pulled) = 1;
    else
        % Empirical estimator of  $\|x\|_2$ , used for ridge
regression.
        penalty = norm(lsX)/size(lsX,1);
        hat_beta = (lsX'*lsX + penalty^2 * eye(d)) \ (lsX' * lsY);
        betahat(arm_pulled,:) = hat_beta';
    end
else
    [XtopX_inv(:, :, arm_pulled), betahat_vertical] =
rankoneupdate(...
    XtopX_inv(:, :, arm_pulled), betahat(arm_pulled, :)',
x, ...
    reward_vector(t, arm_pulled));
    betahat(arm_pulled,:) = betahat_vertical';
end

%----- Fourth, update lambda_min and check whether we should
%switch
if (t==t0)
    for i=1:k
        if(all(XtopX_inv(:, :, i))==0)
            lam_hat(i) = 0;
        else
            lam_hat(i) = 1 / max(eig(XtopX_inv(:, :, i)));
        end
    end
end

```

```

    lam0 = 0.5 * min(lam_hat) / t0;

    if (lam0 < min_eig_threshold)
        switch_to_olsbandit = 1;
        switched = t;
        if(verbose == 1)
            fprintf(...
                'GF: switched at t=t0=%d, lam_hat vector=%s. \n', ...
                t, sprintf('%d ', lam_hat) ...
            );
        end
    end
elseif (t>(2*t0))
    lam_hat(arm_pulled) = 1/max(eig(XtopX_inv(:, :, arm_pulled)));
    [lmmmin, bad_arm]=min(lam_hat);
    if (lmmmin < (lam0*t/4));
        switch_to_olsbandit = 1;
        switched = t;
        if(verbose == 1)
            fprintf('GF: switched at t=%d, lam_hat vector=%s,
lambdamin/t=%f, lam0/4=%f, bad_arm=%d. \n', ...
                t, sprintf('%d ', lam_hat), lmmmin/t, lam0/4,
bad_arm);
        end
    end
end
if(verbose==1)
    if (mod(t,500)==0)
        if(noise_input==1)
            fprintf('GF: t=%d, parameter estimation error = %f.\n',
t, ...
                norm(b - betahat, 'fro'));
        else
            fprintf('GF: t=%d, regret = %f.\n', t, ...
                regret(t));
        end
    end
end
t = t + 1;
end
% If t<T it means that the algorithm has switch to OLS bandit before
the
% time horizon finishes. We need to execute OLS bandit for the
remaining
% time periods.
if(t<=T)
    if(verbose==1)
        fprintf(...
            'GF: Running OLS bandit for the remaining %d periods.\n',
T-t+1);
    end
    if(noise_input==1)
        [regret_OLS, ~, pull_ind_OLS] = runOLSbandit(...)

```

```

        k, T-t+1, d, ...
        b, sigma_e, sigma_x, xmax, h, ...
        q, verbose, X(t:T,:), noise_input, e(t:T));
    regret(t:T) = regret(t-1) + regret_OLS;
    pull_ind(t:T, :) = pull_ind_OLS;

else
    % This version will be mainly used when working with real
    datasets.
    % The main idea for improving the regret of Greedy-First
    (which is
    % quite effective when dimension d is large) is to reuse the
    first
    % random_initialization*k samples gathered at the beginning of
    % compiling greedy algorithm. As these samples are random
    samples,
    % they can be safely used as the initialization (or forced
    % sampling) samples of OLS. The way we do it here is that we
    feed
    % these samples again to the OLS, while we deduct the regret
    of
    % these first random_initialization*k rounds, when Greedy-
    First
    % switches to OLS Bandit.

    ind = [1:random_initialization * k, t:T];
    X_aug = X(ind, :);
    rewards_aug = rewards(ind, :);
    [regret_OLS, ~, pull_ind_OLS] = runOLSbandit(...
        k, length(ind), d, ...
        b, sigma_e, sigma_x, xmax, h, ...
        q, verbose, X_aug, noise_input, rewards_aug);
    regret(t:T) = regret(t-1) + ...
        regret_OLS((random_initialization * k + 1):end) - ...
        regret_OLS(random_initialization*k);
    pull_ind(t:T, :) = pull_ind_OLS(...
        (random_initialization * k + 1):end, :);
end

end

fractions = mean(pull_ind); %fraction of times each arm is pulled
if(verbose==1)
    fprintf('GF: Fraction of pulls = %f. \n', fractions)
    fprintf('GF: Total regret occured = %f. \n', regret(end))
end

end

```

Published with MATLAB® R2015a